

```
In [103]: # Supress Warnings
import warnings
warnings.filterwarnings('ignore')

In [104]: # Importing all required packages

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
from sklearn.preprocessing import scale
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge, Lasso
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import GridSearchCV

# Importing RFE and LinearRegression
from sklearn.feature_selection import RFE
from sklearn.linear_model import LinearRegression
```

Step 1: Reading and Understanding the Data

```
In [105]: df = pd.read_csv("auto-mpg-1.csv")
```

```
In [106]: df.head()
```

```
Out[106]:
```

	name	year	selling_price	km_driven	fuel	seller_type	transmission	owner	mileage	engine	max_power	torque	seats
0	Maruti Swift Dzire VDI	2014	450000	145500	Diesel	Individual	Manual	First Owner	23.4 kmpl	1248 CC	74 bhp	190Nm@ 2000rpm	5.0
1	Skoda Rapid 1.5 TDI Ambition	2014	370000	120000	Diesel	Individual	Manual	Second Owner	21.14 kmpl	1498 CC	103.52 bhp	250Nm@ 1500-2500rpm	5.0
2	Honda City 2017-2020 EXi	2006	158000	140000	Petrol	Individual	Manual	Third Owner	17.7 kmpl	1497 CC	78 bhp	12.7@ 2,700(kgm@ rpm)	5.0
3	Hyundai i20 Sportz Diesel	2010	225000	127000	Diesel	Individual	Manual	First Owner	23.0 kmpl	1396 CC	90 bhp	22.4 kgm at 1750-2750rpm	5.0
4	Maruti Swift VXi BSIII	2007	130000	120000	Petrol	Individual	Manual	First Owner	16.1 kmpl	1298 CC	88.2 bhp	11.5@ 4,500(kgm@ rpm)	5.0

```
In [107]: df = df.drop(['torque'], axis = 1)
```

```
In [108]: print("***** Info *****")
print(df.info())
print("***** Shape *****")
print(df.shape)
print("***** Columns having null values *****")
print(df.isnull().any())
print("***** Describe *****")
df.describe()
```

```
***** Info *****
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8128 entries, 0 to 8127
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   name             8128 non-null   object
1   year             8128 non-null   int64
2   selling_price    8128 non-null   int64
3   km_driven        8128 non-null   int64
4   fuel             8128 non-null   object
5   seller_type      8128 non-null   object
6   transmission     8128 non-null   object
7   owner            8128 non-null   object
8   mileage          7907 non-null   object
9   engine           7907 non-null   object
10  max_power        7913 non-null   object
11  seats            7907 non-null   float64
dtypes: float64(1), int64(3), object(8)
memory usage: 762.1+ KB
None
***** Shape *****
(8128, 12)
***** Columns having null values *****
name             False
year             False
selling_price     False
km_driven        False
fuel             False
seller_type      False
transmission     False
owner            False
mileage          True
engine           True
max_power        True
seats            True
dtype: bool
***** Describe *****
```

```
Out[108]:
```

	year	selling_price	km_driven	seats
count	8128.000000	8.128000e+03	8.128000e+03	7907.000000
mean	2013.804011	6.382718e+05	6.981951e+04	5.416719
std	4.044249	8.062534e+05	5.655055e+04	0.959588
min	1983.000000	2.999900e+04	1.000000e+00	2.000000
25%	2011.000000	2.549990e+05	3.500000e+04	5.000000
50%	2015.000000	4.500000e+05	6.000000e+04	5.000000
75%	2017.000000	6.750000e+05	9.800000e+04	5.000000
max	2020.000000	1.000000e+07	2.360457e+06	14.000000

```
In [109]: df.select_dtypes(include='object').isnull().sum()[df.select_dtypes(include='object').isnull().sum()>0]
```

```
Out[109]: mileage      221
engine      221
max_power   215
dtype: int64
```

```
In [110]: df = df.dropna()
```

```
In [111]: df.select_dtypes(include='object').isnull().sum()[df.select_dtypes(include='object').isnull().sum()>0]
```

```
Out[111]: Series([], dtype: int64)
```

No more null values

```
In [112]: df['mileage'] = df['mileage'].str.replace('kmpl', '')
df['mileage'] = df['mileage'].str.replace('km/kg', '')
df['engine'] = df['engine'].str.replace('CC', '')
df['max_power'] = df['max_power'].str.replace('bhp', '')
```

```
In [113]: #deleting rows that had a unit but no numeric values and converting all values to type float
df['mileage'] = pd.to_numeric(df['mileage'],errors = 'coerce')
df['engine'] = pd.to_numeric(df['engine'],errors = 'coerce')
df['max_power'] = pd.to_numeric(df['max_power'],errors = 'coerce')
```

```
In [114]: df = df.dropna()
```

```
In [115]: df
```

```
Out[115]:
```

	name	year	selling_price	km_driven	fuel	seller_type	transmission	owner	mileage	engine	max_power	seats
0	Maruti Swift Dzire VDI	2014	450000	145500	Diesel	Individual	Manual	First Owner	23.40	1248	74.00	5.0
1	Skoda Rapid 1.5 TDI Ambition	2014	370000	120000	Diesel	Individual	Manual	Second Owner	21.14	1498	103.52	5.0
2	Honda City 2017-2020 EXi	2006	158000	140000	Petrol	Individual	Manual	Third Owner	17.70	1497	78.00	5.0
3	Hyundai i20 Sportz Diesel	2010	225000	127000	Diesel	Individual	Manual	First Owner	23.00	1396	90.00	5.0
4	Maruti Swift VXi BSIII	2007	130000	120000	Petrol	Individual	Manual	First Owner	16.10	1298	88.20	5.0
...
8123	Hyundai i20 Magna	2013	320000	110000	Petrol	Individual	Manual	First Owner	18.50	1197	82.85	5.0
8124	Hyundai Verna CRDi SX	2007	135000	119000	Diesel	Individual	Manual	Fourth & Above Owner	16.80	1493	110.00	5.0
8125	Maruti Swift Dzire ZDi	2009	382000	120000	Diesel	Individual	Manual	First Owner	19.30	1248	73.90	5.0
8126	Tata Indigo CR4	2013	290000	25000	Diesel	Individual	Manual	First Owner	23.57	1396	70.00	5.0
8127	Tata Indigo CR4	2013	290000	25000	Diesel	Individual	Manual	First Owner	23.57	1396	70.00	5.0

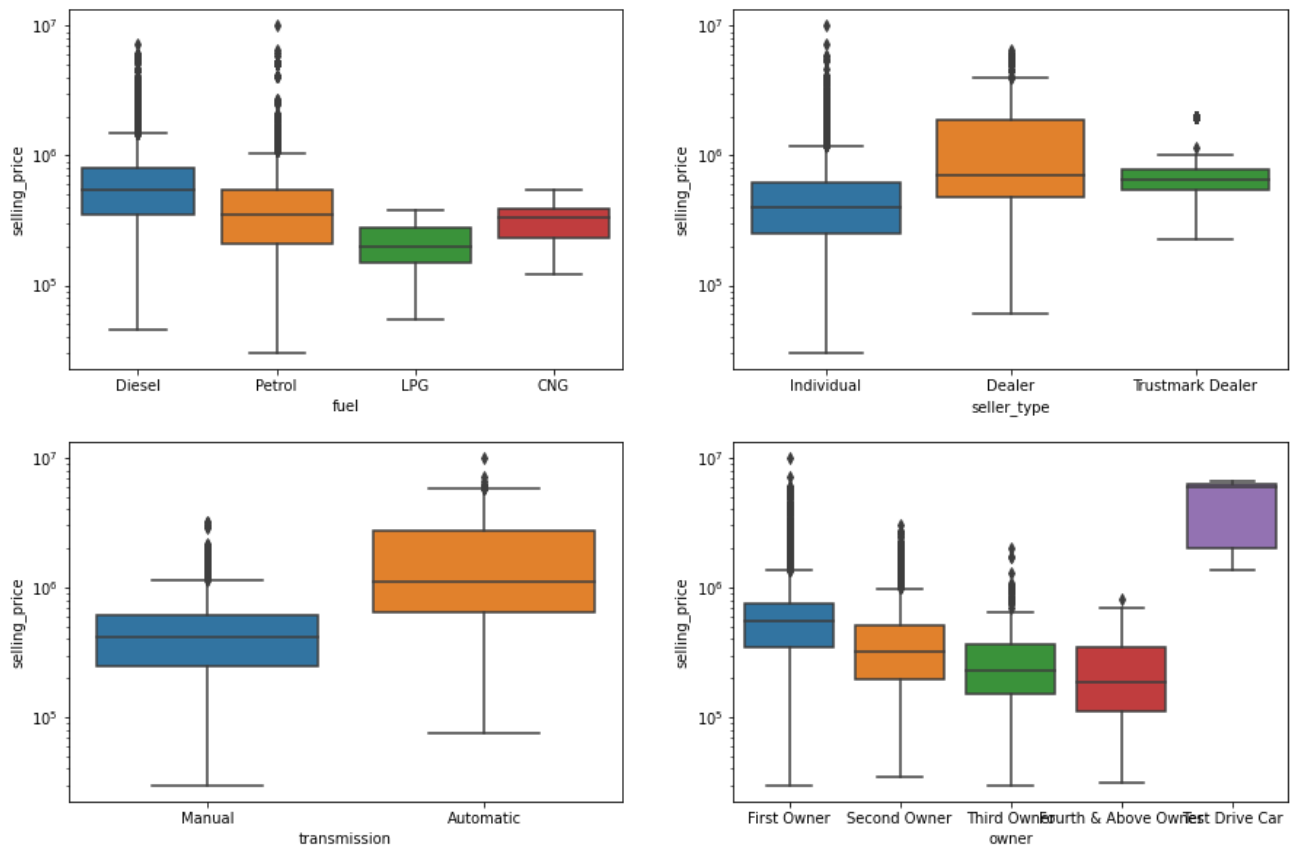
7906 rows × 12 columns

```
In [116]: x = df.drop(columns=['selling_price'])
y = df['selling_price']
```

Plotting the relationships between the label (Selling Price) and the discrete features (fuel type, Seller type, transmission) using a small multiple of box plots

```
In [117]: disc_features = ['fuel', 'seller_type', 'transmission', 'owner']
df1 = X[disc_features]
target = y
```

```
In [118]: fig, ax = plt.subplots(2,2,figsize=(15,10))
for var, subplot in zip(disc_features, ax.flatten()):
    plot = sns.boxplot(x=var, y=target, data = df1, ax=subplot)
    plot.set_yscale("log")
```



From the plots for the continuous variable, we observe a linear relationship as stated earlier. Additionally, the discrete features also can potentially be modeled in a linear fashion. Hence, linear regressions seem to be a suitable model for the given problem.

```
In [119]: # Check the correlation of numerical columns
```

```
plt.figure(figsize = (20, 10))  
sns.heatmap(df.corr(), annot = True, cmap="Greens")  
plt.show()
```



In [120]: X

Out[120]:

	name	year	km_driven	fuel	seller_type	transmission	owner	mileage	engine	max_power	seats
0	Maruti Swift Dzire VDI	2014	145500	Diesel	Individual	Manual	First Owner	23.40	1248	74.00	5.0
1	Skoda Rapid 1.5 TDI Ambition	2014	120000	Diesel	Individual	Manual	Second Owner	21.14	1498	103.52	5.0
2	Honda City 2017-2020 EXi	2006	140000	Petrol	Individual	Manual	Third Owner	17.70	1497	78.00	5.0
3	Hyundai i20 Sportz Diesel	2010	127000	Diesel	Individual	Manual	First Owner	23.00	1396	90.00	5.0
4	Maruti Swift VXI BSIII	2007	120000	Petrol	Individual	Manual	First Owner	16.10	1298	88.20	5.0
...
8123	Hyundai i20 Magna	2013	110000	Petrol	Individual	Manual	First Owner	18.50	1197	82.85	5.0
8124	Hyundai Verna CRDi SX	2007	119000	Diesel	Individual	Manual	Fourth & Above Owner	16.80	1493	110.00	5.0
8125	Maruti Swift Dzire ZDi	2009	120000	Diesel	Individual	Manual	First Owner	19.30	1248	73.90	5.0
8126	Tata Indigo CR4	2013	25000	Diesel	Individual	Manual	First Owner	23.57	1396	70.00	5.0
8127	Tata Indigo CR4	2013	25000	Diesel	Individual	Manual	First Owner	23.57	1396	70.00	5.0

7906 rows x 11 columns

In [121]: X['year'] = 2020 - X['year']

In [122]: *#dropping the car name as it is irrelevant.*
X.drop(["name"],axis = 1,inplace=True)

#check out the dataset with new changes
X.head()

Out[122]:

	year	km_driven	fuel	seller_type	transmission	owner	mileage	engine	max_power	seats
0	6	145500	Diesel	Individual	Manual	First Owner	23.40	1248	74.00	5.0
1	6	120000	Diesel	Individual	Manual	Second Owner	21.14	1498	103.52	5.0
2	14	140000	Petrol	Individual	Manual	Third Owner	17.70	1497	78.00	5.0
3	10	127000	Diesel	Individual	Manual	First Owner	23.00	1396	90.00	5.0
4	13	120000	Petrol	Individual	Manual	First Owner	16.10	1298	88.20	5.0

In [123]: fuel_ohc = pd.get_dummies(X['fuel'])
seller_type_ohc = pd.get_dummies(X['seller_type'])
transmission_ohc = pd.get_dummies(X['transmission'])
owner_ohc = pd.get_dummies(X['owner'])

X = pd.concat([X, fuel_ohc,seller_type_ohc,transmission_ohc,owner_ohc], axis=1)
X = X.drop(columns=['seller_type','fuel','owner','transmission'])

In [124]: *#dataframe after one hot encoding for categorical variables*
X.head()

Out[124]:

	year	km_driven	mileage	engine	max_power	seats	CNG	Diesel	LPG	Petrol	Dealer	Individual	Trustmark Dealer	Automatic	Manual	Fi Owr
0	6	145500	23.40	1248	74.00	5.0	0	1	0	0	0	1	0	0	1	
1	6	120000	21.14	1498	103.52	5.0	0	1	0	0	0	1	0	0	1	
2	14	140000	17.70	1497	78.00	5.0	0	0	0	1	0	1	0	0	1	
3	10	127000	23.00	1396	90.00	5.0	0	1	0	0	0	1	0	0	1	
4	13	120000	16.10	1298	88.20	5.0	0	0	0	1	0	1	0	0	1	

```
In [125]: y
```

```
Out[125]: 0      450000
          1      370000
          2      158000
          3      225000
          4      130000
          ...
          8123   320000
          8124   135000
          8125   382000
          8126   290000
          8127   290000
          Name: selling_price, Length: 7906, dtype: int64
```

Train-test-split

```
In [126]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [127]: lm = LinearRegression()
          lm.fit(X_train, y_train)

          # running RFE
          rfe = RFE(lm, 10)
          rfe = rfe.fit(X_train, y_train)
```

```
In [128]: col = X_train.columns[rfe.support_]

          # assign the 10 features selected using RFE to a dataframe and view them

          temp_df = pd.DataFrame(list(zip(X_train.columns, rfe.support_, rfe.ranking_)), columns=['Variable', 'rfe_s
          temp_df = temp_df.loc[temp_df['rfe_support'] == True]
          temp_df.reset_index(drop=True, inplace=True)

          temp_df
```

```
Out[128]:
```

	Variable	rfe_support	rfe_ranking
0	Dealer	True	1
1	Individual	True	1
2	Trustmark Dealer	True	1
3	Automatic	True	1
4	Manual	True	1
5	First Owner	True	1
6	Fourth & Above Owner	True	1
7	Second Owner	True	1
8	Test Drive Car	True	1
9	Third Owner	True	1

```
In [129]: # Assign the 10 columns to X_train_rfe

          X_train_rfe = X_train[col]
```

```
In [130]: # Associate the new 10 columns to X_train and X_test for further analysis

          X_train = X_train_rfe[X_train_rfe.columns]
          X_test = X_test[X_train.columns]
```

Model Building and Evaluation

```

In [131]: # list pf alphas

params = {'alpha': [0.0001, 0.001, 0.01, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0, 4.0,
                    9.0, 10.0, 20, 50, 100, 500, 1000 ]}

ridge = Ridge()

# cross validation

folds = 5
ridge_model_cv = GridSearchCV(estimator = ridge,
                              param_grid = params,
                              scoring= 'neg_mean_absolute_error',
                              cv = folds,
                              return_train_score=True,
                              verbose = 1)
ridge_model_cv.fit(X_train, y_train)

```

Fitting 5 folds for each of 27 candidates, totalling 135 fits

```

Out[131]: GridSearchCV(cv=5, estimator=Ridge(),
                      param_grid={'alpha': [0.0001, 0.001, 0.01, 0.1, 0.2, 0.3, 0.4, 0.5,
                                             0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0, 4.0, 5.0,
                                             6.0, 7.0, 8.0, 9.0, 10.0, 20, 50, 100, 500,
                                             1000]},
                      return_train_score=True, scoring='neg_mean_absolute_error',
                      verbose=1)

```



```
In [132]: # display the mean scores
```

```
ridge_cv_results = pd.DataFrame(ridge_model_cv.cv_results_)
ridge_cv_results = ridge_cv_results[ridge_cv_results['param_alpha']<=500]
ridge_cv_results[['param_alpha', 'mean_train_score', 'mean_test_score', 'rank_test_score']].sort_values()
```

```
Out[132]:
```


	param_alpha	mean_train_score	mean_test_score	rank_test_score
25	500	-351943.427660	-352149.911672	1
24	100	-351944.111765	-352223.388635	2
5	0.3	-352519.917714	-352942.440142	3
4	0.2	-352477.818126	-352943.355404	4
3	0.1	-352433.127534	-352952.400830	5
2	0.01	-352401.019199	-352961.852665	6
1	0.001	-352397.879777	-352962.874129	7
0	0.0001	-352397.564952	-352962.977083	8
6	0.4	-352559.656170	-352978.732596	9
7	0.5	-352597.257289	-353013.466791	10
8	0.6	-352632.842121	-353046.388068	11
9	0.7	-352666.582104	-353077.686892	12
10	0.8	-352698.742403	-353107.585754	13
23	50	-352834.013288	-353122.281073	14
11	0.9	-352729.358281	-353136.115526	15
12	1.0	-352758.495660	-353163.348321	16
13	2.0	-352987.839185	-353377.488084	17
14	3.0	-353139.774391	-353517.935661	18
15	4.0	-353244.114353	-353613.767038	19
16	5.0	-353317.188349	-353680.566119	20
22	20	-353384.213679	-353701.776255	21
17	6.0	-353368.595657	-353727.069423	22
18	7.0	-353404.217586	-353758.430509	23
19	8.0	-353428.808274	-353778.624775	24
20	9.0	-353445.042192	-353790.353159	25
21	10.0	-353454.233247	-353795.603748	26

```
In [133]: # plotting mean test and train scoes with alpha

ridge_cv_results['param_alpha'] = ridge_cv_results['param_alpha'].astype('int32')

# plotting

plt.plot(ridge_cv_results['param_alpha'], ridge_cv_results['mean_train_score'])
plt.plot(ridge_cv_results['param_alpha'], ridge_cv_results['mean_test_score'])
plt.xlabel('alpha')
plt.ylabel('Negative Mean Absolute Error')
plt.title("Negative Mean Absolute Error and alpha")
plt.legend(['train score', 'test score'], loc='upper right')
plt.show()
```



```
In [134]: # get the best estimator for lambda

ridge_model_cv.best_estimator_
```

Out[134]: Ridge(alpha=500)

```
In [135]: # check the coefficient values with lambda = 10

alpha = 500
ridge = Ridge(alpha=alpha)

ridge.fit(X_train, y_train)
ridge.coef_
```

Out[135]: array([238767.39832227, -175765.31695917, -63002.0813631 ,
449820.57748154, -449820.57748154, 138455.66268397,
-40899.26938719, -58164.510249 , 27902.33582309,
-67294.21887086])

```
In [136]: # Put the Features and coefficient in a dataframe

ridge_df = pd.DataFrame({'Features':X_train.columns, 'Coefficient':ridge.coef_.round(4)})
ridge_df.reset_index(drop=True, inplace=True)
ridge_df
```

Out[136]:

	Features	Coefficient
0	Dealer	238767.3983
1	Individual	-175765.3170
2	Trustmark Dealer	-63002.0814
3	Automatic	449820.5775
4	Manual	-449820.5775
5	First Owner	138455.6627
6	Fourth & Above Owner	-40899.2694
7	Second Owner	-58164.5102
8	Test Drive Car	27902.3358
9	Third Owner	-67294.2189

```
In [137]: # Assign the Features and their coefficient values to a dictionary which would be used while plotting the
```

```
ridge_coeff_dict = dict(pd.Series(ridge.coef_.round(4), index = X_train.columns))
ridge_coeff_dict
```

```
Out[137]: {'Dealer': 238767.3983,
'Individual': -175765.317,
'Trustmark Dealer': -63002.0814,
'Automatic': 449820.5775,
'Manual': -449820.5775,
'First Owner': 138455.6627,
'Fourth & Above Owner': -40899.2694,
'Second Owner': -58164.5102,
'Test Drive Car': 27902.3358,
'Third Owner': -67294.2189}
```

```
In [138]: # Do an RFE to minimise the features to 8
X_train_ridge = X_train[ridge_df.Features]
```

```
lm = LinearRegression()
lm.fit(X_train_ridge, y_train)

# running RFE
rfe = RFE(lm, 8)
rfe = rfe.fit(X_train_ridge, y_train)
```

```
In [139]: # Method to get the coefficient values
```

```
def find(x):
    return ridge_coeff_dict[x]

# Assign top 10 features to a temp dataframe for further display in the bar plot

templ_df = pd.DataFrame(list(zip( X_train_ridge.columns, rfe.support_, rfe.ranking_)), columns=['Features', 'rfe_support', 'rfe_ranking'])
templ_df = templ_df.loc[templ_df['rfe_support'] == True]
templ_df.reset_index(drop=True, inplace=True)

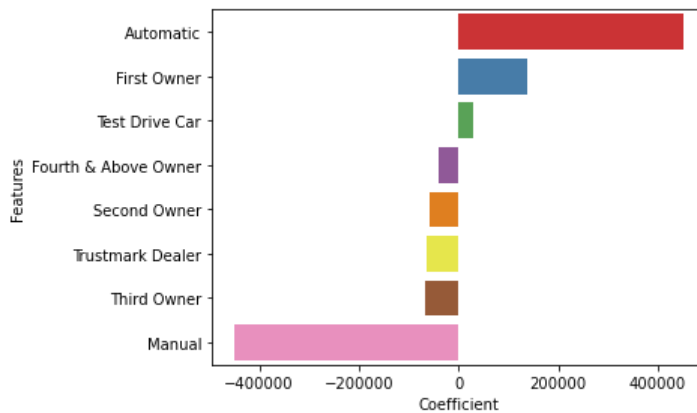
templ_df['Coefficient'] = templ_df['Features'].apply(find)
templ_df = templ_df.sort_values(by=['Coefficient'], ascending=False)
templ_df = templ_df.head(10)
templ_df
```

```
Out[139]:
```

	Features	rfe_support	rfe_ranking	Coefficient
1	Automatic	True	1	449820.5775
3	First Owner	True	1	138455.6627
6	Test Drive Car	True	1	27902.3358
4	Fourth & Above Owner	True	1	-40899.2694
5	Second Owner	True	1	-58164.5102
0	Trustmark Dealer	True	1	-63002.0814
7	Third Owner	True	1	-67294.2189
2	Manual	True	1	-449820.5775

```
In [140]: # bar plot to determine the variables that would affect pricing most using ridge regression
```

```
plt.figure(figsize=(20,20))
plt.subplot(4,3,1)
sns.barplot(y = 'Features', x='Coefficient', palette='Set1', data = templ_df)
plt.show()
```



Lasso

```
In [141]: lasso = Lasso()

# list of alphas

params = {'alpha': [0.0001, 0.0002, 0.0003, 0.0004, 0.0005, 0.001, 0.002, 0.003, 0.004, 0.005, 0.01]}

# cross validation

folds = 5
lasso_model_cv = GridSearchCV(estimator = lasso,
                              param_grid = params,
                              scoring= 'neg_mean_absolute_error',
                              cv = folds,
                              return_train_score=True,
                              verbose = 1)

lasso_model_cv.fit(X_train, y_train)
```

Fitting 5 folds for each of 11 candidates, totalling 55 fits

```
Out[141]: GridSearchCV(cv=5, estimator=Lasso(),
                      param_grid={'alpha': [0.0001, 0.0002, 0.0003, 0.0004, 0.0005,
                                             0.001, 0.002, 0.003, 0.004, 0.005, 0.01]},
                      return_train_score=True, scoring='neg_mean_absolute_error',
                      verbose=1)
```

```
In [142]: # display the mean scores

lasso_cv_results = pd.DataFrame(lasso_model_cv.cv_results_)
lasso_cv_results[['param_alpha', 'mean_train_score', 'mean_test_score', 'rank_test_score']].sort_values()
```

Out[142]:

	param_alpha	mean_train_score	mean_test_score	rank_test_score
10	0.01	-352397.533197	-352962.985755	1
9	0.005	-352397.531579	-352962.987143	2
8	0.004	-352397.531256	-352962.987420	3
7	0.003	-352397.530932	-352962.987698	4
6	0.002	-352397.530609	-352962.987976	5
5	0.001	-352397.530285	-352962.988253	6
4	0.0005	-352397.530123	-352962.988392	7
3	0.0004	-352397.530091	-352962.988420	8
2	0.0003	-352397.530059	-352962.988448	9
1	0.0002	-352397.530026	-352962.988475	10
0	0.0001	-352397.529994	-352962.988503	11

```
In [143]: lasso_cv_results['param_alpha'] = lasso_cv_results['param_alpha'].astype('float64')
```

```
In [144]: # get the best estimator for lambda

lasso_model_cv.best_estimator_
```

Out[144]: Lasso(alpha=0.01)

```
In [145]: alpha = 0.01

lasso = Lasso(alpha=alpha)

lasso.fit(X_train, y_train)
lasso.coef_
```

Out[145]: array([468827.51521879, 29696.28429588, -119603.5397049 ,
1203787.83059786, -0. , 217243.44913252,
-112922.03676589, 15582.03475208, 3371031.66951676,
-52946.59414298])

```
In [146]: # Put the shortlisted Features and coefficient in a dataframe

lasso_df = pd.DataFrame({'Features':X_train.columns, 'Coefficient':lasso.coef_.round(4)})
lasso_df = lasso_df[lasso_df['Coefficient'] != 0.00]
lasso_df.reset_index(drop=True, inplace=True)
lasso_df
```

Out[146]:

	Features	Coefficient
0	Dealer	4.688275e+05
1	Individual	2.969628e+04
2	Trustmark Dealer	-1.196035e+05
3	Automatic	1.203788e+06
4	First Owner	2.172434e+05
5	Fourth & Above Owner	-1.129220e+05
6	Second Owner	1.558203e+04
7	Test Drive Car	3.371032e+06
8	Third Owner	-5.294659e+04

```
In [147]: # Put the Features and Coefficients in dictionary

lasso_coeff_dict = dict(pd.Series(lasso.coef_, index = X_train.columns))
lasso_coeff_dict

Out[147]: {'Dealer': 468827.51521878934,
'Individual': 29696.284295879646,
'Trustmark Dealer': -119603.53970490451,
'Automatic': 1203787.8305978617,
'Manual': -0.0,
'First Owner': 217243.44913252094,
'Fourth & Above Owner': -112922.03676589063,
'Second Owner': 15582.034752076945,
'Test Drive Car': 3371031.6695167557,
'Third Owner': -52946.594142984024}

In [148]: # Do an RFE to minimise the features to 8

X_train_lasso = X_train[lasso_df.Features]

lm = LinearRegression()
lm.fit(X_train_lasso, y_train)

# running RFE

rfe = RFE(lm, 8)
rfe = rfe.fit(X_train_lasso, y_train)

In [149]: # Method to get the coefficient values

def find(x):
    return lasso_coeff_dict[x]

# Assign top 8 features to a temp dataframe for further display in the bar plot

temp2_df = pd.DataFrame(list(zip( X_train_lasso.columns, rfe.support_, rfe.ranking_)), columns=['Features', 'rfe_support', 'rfe_ranking'])
temp2_df = temp2_df.loc[temp2_df['rfe_support'] == True]
temp2_df.reset_index(drop=True, inplace=True)


temp2_df['Coefficient'] = temp2_df['Features'].apply(find)
temp2_df = temp2_df.sort_values(by=['Coefficient'], ascending=False)
temp2_df = temp2_df.head(10)
temp2_df
```

Out[149]:

	Features	rfe_support	rfe_ranking	Coefficient
6	Test Drive Car	True	1	3.371032e+06
0	Dealer	True	1	4.688275e+05
3	First Owner	True	1	2.172434e+05
1	Individual	True	1	2.969628e+04
5	Second Owner	True	1	1.558203e+04
7	Third Owner	True	1	-5.294659e+04
4	Fourth & Above Owner	True	1	-1.129220e+05
2	Trustmark Dealer	True	1	-1.196035e+05

```
In [150]: # bar plot to determine the variables that would affect pricing most using ridge regression
```

```
plt.figure(figsize=(20,20))
plt.subplot(4,3,1)
sns.barplot(y = 'Features', x='Coefficient', palette='Set1', data = temp2_df)
plt.show()
```



```
In [151]: # Check the mean squared error
mean_squared_error(y_test, ridge.predict(X_test))
```

```
Out[151]: 417808486827.479
```

```
In [152]: # Check the mean squared error
mean_squared_error(y_test, lasso.predict(X_test))
```

```
Out[152]: 403341306275.08484
```

Conclusion :

- The optimal lambda value in case of Ridge and Lasso is as below:
 - Ridge - 10
 - Lasso - 0.01
- The mean squared error for the dataset is really high. Hence, linear regression may not be a model to apply on this dataset.
- The mean squared error of Lasso is slightly lower than that of Ridge
- Therefore, the variables predicted by Lasso in the above bar chart are significant variables for predicting the price of the car