

In [1]:

```
import numpy as np
import matplotlib.pyplot as plt
import random
from collections import defaultdict
from aima.mdp4e import MDP, policy_evaluation

class PassiveDUEAgent:
    """
    Passive (non-learning) agent that uses direct utility estimation
    on a given MDP and policy.

    import sys
    from mdp import sequential_decision_environment
    north = (0, 1)
    south = (0,-1)
    west = (-1, 0)
    east = (1, 0)
    policy = {(0, 2): east, (1, 2): east, (2, 2): east, (3, 2): None, (0, 1): north, (2, 1):
              (3, 1): None, (0, 0): north, (1, 0): west, (2, 0): west, (3, 0): west,}
    agent = PassiveDUEAgent(policy, sequential_decision_environment)
    for i in range(200):
        run_single_trial(agent,sequential_decision_environment)
        agent.estimate_U()
    agent.U[(0, 0)] > 0.2
    True
    """

    def __init__(self, pi, mdp):
        self.pi = pi
        self.mdp = mdp
        self.U = {}
        self.s = None
        self.a = None
        self.s_history = []
        self.r_history = []
        self.init = mdp.init

    def __call__(self, percept):
        s1, r1 = percept
        self.s_history.append(s1)
        self.r_history.append(r1)
        ##
        ##
        if s1 in self.mdp.terminals:
            self.s = self.a = None
        else:
            self.s, self.a = s1, self.pi[s1]
        return self.a

    def estimate_U(self):
        # this function can be called only if the MDP has reached a terminal state
        # it will also reset the mdp history
        assert self.a is None, 'MDP is not in terminal state'
        assert len(self.s_history) == len(self.r_history)
        # calculating the utilities based on the current iteration
        U2 = {s: [] for s in set(self.s_history)}
        for i in range(len(self.s_history)):
            s = self.s_history[i]
            U2[s] += [sum(self.r_history[i:])]
        U2 = {k: sum(v) / max(len(v), 1) for k, v in U2.items()}
        # resetting history
        self.s_history, self.r_history = [], []
```

```

# setting the new utilities to the average of the previous
# iteration and this one
for k in U2.keys():
    if k in self.U.keys():
        self.U[k] = (self.U[k] + U2[k]) / 2
    else:
        self.U[k] = U2[k]
return self.U

def update_state(self, percept):
    """To be overridden in most cases. The default case
    assumes the percept to be of type (state, reward)"""
    return percept

def direct_utility(self, state, trans, utility, reward, gamma):
    """
    function: direct_utility
    params: vevtor ints state, vevtor ints transition, vector ints utility, float reward
    does: calculates the bellman equation for the utility of a state
    returns: the utility of a state float
    """
    actions = [0.0, 0.0, 0.0, 0.0]
    for action in range(0, 4):
        actions[action] = np.sum(np.multiply(utility, np.dot(state, trans[:, :, action])))
    out = reward + gamma * np.max(actions)
    return out

class PassiveADPAgent:
    """
    Passive (non-learning) agent that uses adaptive dynamic programming
    on a given MDP and policy.

    import sys
    from mdp import sequential_decision_environment
    north = (0, 1)
    south = (0, -1)
    west = (-1, 0)
    east = (1, 0)
    policy = {(0, 2): east, (1, 2): east, (2, 2): east, (3, 2): None, (0, 1): north, (2, 1):
              (3, 1): None, (0, 0): north, (1, 0): west, (2, 0): west, (3, 0): west,}
    agent = PassiveADPAgent(policy, sequential_decision_environment)
    for i in range(100):
        run_single_trial(agent, sequential_decision_environment)

    agent.U[(0, 0)] > 0.2
    True
    agent.U[(0, 1)] > 0.2
    True
    """

class ModelMDP(MDP):
    """Class for implementing modified Version of input MDP with
    an editable transition model P and a custom function T."""

    def __init__(self, init, actlist, terminals, gamma, states):
        super().__init__(init, actlist, terminals, states=states, gamma=gamma)
        nested_dict = lambda: defaultdict(nested_dict)
        # StackOverflow:whats-the-best-way-to-initialize-a-dict-of-dicts-in-python
        self.P = nested_dict()

    def T(self, s, a):
        """Return a list of tuples with probabilities for states
        based on the learnt model P."""
        return [(prob, res) for (res, prob) in self.P[(s, a)].items()]

```

```

def __init__(self, pi, mdp):
    self.pi = pi
    self.mdp = PassiveADPAgent.ModelMDP(mdp.init, mdp.actlist,
                                          mdp.terminals, mdp.gamma, mdp.states)

    self.U = {}
    self.Nsa = defaultdict(int)
    self.Nsl_sa = defaultdict(int)
    self.s = None
    self.a = None
    self.visited = set() # keeping track of visited states

def __call__(self, percept):
    s1, r1 = percept
    mdp = self.mdp
    R, P, terminals, pi = mdp.reward, mdp.P, mdp.terminals, self.pi
    s, a, Nsa, Nsl_sa, U = self.s, self.a, self.Nsa, self.Nsl_sa, self.U

    if s1 not in self.visited: # Reward is only known for visited state.
        U[s1] = R[s1] = r1
        self.visited.add(s1)
    if s is not None:
        Nsa[(s, a)] += 1
        Nsl_sa[(s1, s, a)] += 1
        # for each t such that Nsa[t, s, a] is nonzero
        for t in [res for (res, state, act), freq in Nsl_sa.items()
                  if (state, act) == (s, a) and freq != 0]:
            P[(s, a)][t] = Nsl_sa[(t, s, a)] / Nsa[(s, a)]

    self.U = policy_evaluation(pi, U, mdp)
    ##
    ##
    self.Nsa, self.Nsl_sa = Nsa, Nsl_sa
    if s1 in terminals:
        self.s = self.a = None
    else:
        self.s, self.a = s1, self.pi[s1]
    return self.a

def update_state(self, percept):
    """To be overridden in most cases. The default case
    assumes the percept to be of type (state, reward)."""
    return percept

def adp_rewards(self):
    def reward_function(pos, action):
        if pos in goal:
            return pos, 0
        reward = r
        end_pos = np.array(pos) + np.array(action)
        if -1 in end_pos or 4 in end_pos:
            end_pos = pos
        return end_pos, reward

    gamma = 1
    r = -1
    grid_constraint = 4
    goal = [[0, 0], [grid_constraint - 1, grid_constraint - 1]]
    actions = [[1, 0], [0, 1], [0, -1], [-1, 0]]
    iters = 300
    grid = np.zeros((grid_constraint, grid_constraint), dtype=float)
    state = [[state_1, state_2] for state_1 in range(grid_constraint) for state_2 in range(grid_constraint)]

    results = []
    # loop over total number of iterations set to 400, this can be changed until the

```

```

for item in range(100):
    copy_grid = np.copy(grid)
    states = []
    for s in state:
        weightedRewards = 0
        for action in actions:
            end_pos, reward = reward_function(s, action)
            weightedRewards += (1 / len(actions)) * (reward + (gamma * grid[end_pos]))
        states.append(np.abs(copy_grid[s[0], s[1]] - weightedRewards))
        copy_grid[s[0], s[1]] = weightedRewards
    results.append(states)
    grid = copy_grid
    #if item in [0, 50, 100, 150, 200, 250, 300]:
        #print("Iteration:", item)
        #print(grid)

plt.figure()
plt.plot(results)
plt.show()

```

```

class PassiveTDAgent:

```

```

    """

```

The abstract class for a Passive (non-learning) agent that uses temporal differences to learn utility estimates. Override update_state method to convert percept to state and reward. The mdp being provided should be an instance of a subclass of the MDP Class.

```

import sys
from mdp import sequential_decision_environment
north = (0, 1)
south = (0,-1)
west = (-1, 0)
east = (1, 0)
policy = {(0, 2): east, (1, 2): east, (2, 2): east, (3, 2): None, (0, 1): north, (2, 1):
        (3, 1): None, (0, 0): north, (1, 0): west, (2, 0): west, (3, 0): west,}
agent = PassiveTDAgent(policy, sequential_decision_environment, alpha=lambda n: 60./(5+n))
for i in range(200):
    run_single_trial(agent,sequential_decision_environment)

agent.U[(0, 0)] > 0.2
True
agent.U[(0, 1)] > 0.2
True
    """

```

```

def __init__(self, pi, mdp, alpha=None):

```

```

    self.pi = pi
    self.U = {s: 0. for s in mdp.states}
    self.Ns = {s: 0 for s in mdp.states}
    self.s = None
    self.a = None
    self.r = None
    self.gamma = mdp.gamma
    self.terminals = mdp.terminals

```

```

    if alpha:
        self.alpha = alpha
    else:
        self.alpha = lambda n: 1 / (1 + n) # udacity video

```

```

def __call__(self, percept):

```

```

s1, r1 = self.update_state(percept)
pi, U, Ns, s, r = self.pi, self.U, self.Ns, self.s, self.r
alpha, gamma, terminals = self.alpha, self.gamma, self.terminals
if not Ns[s1]:
    U[s1] = r1
if s is not None:
    Ns[s] += 1
    U[s] += alpha(Ns[s]) * (r + gamma * U[s1] - U[s])
if s1 in terminals:
    self.s = self.a = self.r = None
else:
    self.s, self.a, self.r = s1, pi[s1], r1
return self.a

def update_state(self, percept):
    return percept

def td_rewards(self):
    gamma = 0.1
    r = -1
    grid = 4
    alpha = 0.1
    goal = [[0, 0], [grid - 1, grid - 1]]
    actions = [[1, 0], [0, 1], [0, -1], [-1, 0]]
    iters = 10000
    env = np.zeros((grid, grid))
    states = [[state_1, state_2] for state_1 in range(grid) for state_2 in range(grid)]
    deltas = {(i, j): list() for i in range(grid) for j in range(grid)}

    def init_state():
        """
        function: init_states
        params: none
        returns: ints init state
        """
        init_state = random.choice(states[1:-1])
        return init_state

    def get_next_action():
        """
        function: get_next_action
        params: none
        returns: makes a random choice of direction for the agent
        """
        return random.choice(actions)

    def reward_function_two(pos, action):
        """
        function: reward_function
        params: pos ints the initial position, action ints
        returns: the next position ints, and reward float
        """
        if list(pos) in goal:
            return 0, None
        end_state = np.array(pos) + np.array(action)

        if -1 in list(end_state) or grid in list(end_state):
            end_state = pos
        return r, list(end_state)

    for i in range(iters):
        state = init_state()
        while True:
            action = get_next_action()

```

```

rew, end_state = reward_function_two(state, action)

    if end_state is None:
        break
    # value function

    prev_states = env[state[0], state[1]]
    env[state[0], state[1]] += alpha * (
        rew + gamma * env[end_state[0], end_state[1]] - env[state[0], state[1]]
    )
    deltas[state[0], state[1]].append(float(np.abs(prev_states - env[state[0], state[1]])))
    state = end_state
plt.figure(figsize=(10, 5))
all_series = [list(x)[:50] for x in deltas.values()]
for series in all_series:
    plt.plot(series)
plt.show()

def run_single_trial(agent_program, mdp):
    """Execute trial for given agent_program
    and mdp. mdp should be an instance of subclass
    of mdp.MDP """

    def take_single_action(mdp, s, a):
        """
        Select outcome of taking action a
        in state s. Weighted Sampling.
        """
        x = random.uniform(0, 1)
        cumulative_probability = 0.0
        for probability_state in mdp.T(s, a):
            probability, state = probability_state
            cumulative_probability += probability
            if x < cumulative_probability:
                break
        return state

    current_state = mdp.init
    while True:
        current_reward = mdp.R(current_state)
        percept = (current_state, current_reward)
        next_action = agent_program(percept)
        if next_action is None:
            break
        current_state = take_single_action(mdp, current_state, next_action)

# -----

"""
A 4x3 grid environment that presents the agent with a sequential decision problem.
"""

import sys
from aima.mdp import sequential_decision_environment

north = (0, 1)
south = (0, -1)
west = (-1, 0)
east = (1, 0)
policy = {(0, 2): east, (1, 2): east, (2, 2): east, (3, 2): None, (0, 1): north, (2, 1): north, (3, 1): None, (0, 0): north, (1, 0): west, (2, 0): west, (3, 0): west, }
agentADP = PassiveADPAgent(policy, sequential_decision_environment)
agentTD = PassiveTDAgent(policy, sequential_decision_environment)
agentD = PassiveDUEAgent(policy, sequential_decision_environment)
agentADP.adp_rewards()

```

```
agentTD.td_rewards()  
for i in range(200):  
    run_single_trial(agentD, sequential_decision_environment)  
    agentD.estimate_U()
```

Warning: Transition table is empty.

